

Data Mining Viva Preparation Notes

(Unit-wise, Topic-wise Structured Questions and Answers)

Instructions

This document contains short viva questions and answers arranged unit-wise and topic-wise. Existing questions are retained and new questions are added after each unit.

Viva Questions and Answers – Data Mining (Unit-wise)

Unit 1: Introduction to Data Mining

1. What is Data Mining?

Data mining is the process of discovering patterns and useful information from large datasets.

2. What is Knowledge Discovery in Databases (KDD)?

KDD is the overall process of converting raw data into useful knowledge; data mining is one step in it.

3. Give two motivations for data mining.

Huge data generation and the need for decision-making based on patterns.

4. What are supervised techniques?

Techniques that use labelled data, e.g., classification.

5. What are unsupervised techniques?

Techniques that use unlabelled data, e.g., clustering.

6. What are the types of data mining tasks?

Classification, clustering, association rule mining, prediction, summarization.

7. What are applications of data mining?

Fraud detection, market basket analysis, healthcare, recommendation systems.

8. What is data quality?

The degree to which data is accurate, complete, consistent and reliable.

9. What are data measurements?

Types of measurement scales: nominal, ordinal, interval and ratio.

10. Difference between prediction and classification?

Prediction deals with continuous values; classification deals with categorical values.

11. What is the goal of data mining?

To extract meaningful patterns and actionable knowledge from data.

12. What is descriptive data mining?

It focuses on summarizing the general properties or patterns of data.

13. What is predictive data mining?

It predicts unknown or future values using existing data patterns.

14. What is meant by Big Data?

Extremely large datasets that are difficult to manage with traditional tools.

15. State the 3 Vs of Big Data.

Volume, Velocity, and Variety.

16. What is data heterogeneity?

When data comes from diverse sources and formats.

17. What is noise in data?

Irrelevant or random errors that affect data accuracy.

18. What is outlier?

A data point that differs significantly from other observations.

19. What is metadata?

Data that describes other data.

20. What is OLAP?

Online Analytical Processing; used for fast analysis of multidimensional data.

21. Difference between OLAP and OLTP?

OLAP is for analysis; OLTP is for transaction processing.

22. What is a data warehouse?

A centralized database used for reporting and analysis.

23. What is data cleaning?

The process of detecting and correcting errors in data.

24. Define pattern discovery.

Identifying trends, rules, or regularities in data.

Unit 2: Data Pre-Processing

26. What is data preprocessing?

A set of techniques used to clean and prepare raw data for mining.

27. What is data aggregation?

Combining multiple small objects into larger ones.

28. What is sampling?

Selecting a subset of data to represent the entire dataset.

29. What is dimensionality reduction?

Reducing the number of attributes while preserving important information.

30. Name one dimensionality reduction method.

Principal Component Analysis (PCA).

31. What is feature subset selection?

Selecting the best features that contribute most to prediction.

32. What is feature creation?

Creating new features from existing ones.

33. What is variable transformation?

Changing the scale or format of variables, e.g., normalization.

34. What is normalization?

Scaling data to a small range like [0,1].

35. Why is preprocessing important?

It improves the accuracy and performance of data mining models.

36. What is data cleaning?

Removing noise, errors, and inconsistencies from data.

37. What is data integration?

Combining data from multiple sources into a single dataset.

38. What is data reduction?

Reducing data volume while maintaining the integrity of data.

39. What is PCA used for?

To convert correlated variables into uncorrelated principal components.

40. What is discretization?

Converting continuous data into categorical bins.

41. What is binning?

Grouping a number of continuous values into smaller intervals.

42. What is smoothing?

Removing noise by using techniques like bin means or medians.

43. What is Z-score normalization?

Transforming data using mean and standard deviation.

44. What is min–max normalization formula?

$$v' = \frac{v - \min}{\max - \min}.$$

45. What is data inconsistency?

When the same data shows different values across sources.

46. What is missing data imputation?

Filling missing values using mean, median, or model prediction.

47. What is attribute construction?

Creating new attributes from existing ones to improve model accuracy.

48. What is curse of dimensionality?

When high dimensional data becomes sparse and models perform poorly.

49. Why is sampling used?

To reduce computation time and cost while preserving data characteristics.

50. What is random sampling?

Selecting samples where every data item has equal chance of inclusion.

Unit 3: Cluster Analysis

51. What is clustering?

Grouping similar objects such that intra-cluster similarity is high.

52. What is a similarity measure?

A function that determines how close two data points are.

53. What is Euclidean distance?

A straight-line distance between two points used as a similarity measure.

54. Name two types of clusters.

Well-separated clusters and density-based clusters.

55. What is K-means clustering?

A partitional clustering algorithm that divides data into k clusters.

56. What is a centroid in K-means?

The mean point of all items in a cluster.

57. How does K-means work?

It iteratively assigns points to the nearest centroid and updates centroids.

58. What is the main limitation of K-means?

It requires the number of clusters k in advance.

59. What is cluster validation?

Methods used to evaluate the quality of clusters.

60. How to determine the optimal number of clusters?

Using the Elbow method or Silhouette score.

61. What is hierarchical clustering?

A method that builds clusters in a tree-like structure.

62. Difference between agglomerative and divisive clustering?

Agglomerative: bottom-up; Divisive: top-down.

63. **What is dendrogram?**
A tree diagram representing hierarchical clustering.
64. **What is Manhattan distance?**
Distance calculated as sum of absolute differences of coordinates.
65. **What is cosine similarity?**
It measures similarity based on angle between vectors.
66. **What is density-based clustering?**
Clustering based on dense regions, e.g., DBSCAN.
67. **What is DBSCAN?**
A clustering algorithm based on density and neighborhood parameters.
68. **What is minPts in DBSCAN?**
Minimum number of points to form a dense region.
69. **What is eps in DBSCAN?**
Maximum radius to search for neighboring points.
70. **What is a noise point in DBSCAN?**
A point that does not belong to any cluster.
71. **What is Silhouette score?**
A measure of cluster quality from -1 to 1.
72. **Define intra-cluster distance.**
Average distance between points in the same cluster.
73. **Define inter-cluster distance.**
Distance between cluster centroids.

Unit 4: Association Rule Mining

76. **What is association rule mining?**
Finding interesting relations among items in transaction data.
77. **What is a transaction dataset?**
Data where each record contains a list of purchased items.
78. **What is support?**
The proportion of transactions that contain an itemset.
79. **What is confidence?**
Probability that item Y occurs given X occurs.
80. **What is the Apriori algorithm?**
An algorithm to find frequent itemsets using the Apriori principle.
81. **What is the Apriori principle?**
If an itemset is frequent, all of its subsets must also be frequent.

82. **What are frequent itemsets?**

Itemsets whose support is above the minimum threshold.

83. **What is rule generation?**

Creating association rules from frequent itemsets.

84. **Give one application of association rule mining.**

Market basket analysis.

85. **What is lift?**

Lift measures how much more likely X and Y occur together than expected if independent.

86. **Give the formula for lift.**

$$\text{Lift} = \frac{\text{Support}(XY)}{\text{Support}(X) \times \text{Support}(Y)}$$

87. **What is leverage?**

Difference between actual and expected frequency of X and Y occurring together.

88. **What is conviction?**

Measures how strongly X implies Y .

89. **What is support count?**

The number of transactions containing an itemset.

90. **What is a candidate itemset?**

A possible frequent itemset generated during Apriori.

91. **What is pruning in Apriori?**

Removing candidates with infrequent subsets.

92. **What is FP-growth?**

A fast algorithm to find frequent itemsets using FP-tree.

93. **What is FP-tree?**

A compressed representation of transaction data.

94. **What is closed frequent itemset?**

A frequent itemset with no superset having the same support.

Unit 5: Classification

96. **What is classification?**

Assigning items to predefined categories.

97. **What is the Naive Bayes classifier?**

A probabilistic classifier based on Bayes' theorem assuming feature independence.

98. **What is the K-NN classifier?**

It assigns a label based on the majority class of k nearest points.

99. **What is a decision tree?**

A structure where nodes represent tests and leaves represent classes.

100. What is overfitting?

When a model performs well on training data but poorly on test data.

101. What is a confusion matrix?

A table showing actual vs. predicted classifications.

102. What are evaluation metrics?

Metrics like accuracy, precision, recall, F1-score.

103. What is model evaluation?

Assessing a classifier's performance.

104. What is training data?

Data used to teach the model.

105. What is test data?

Data used to evaluate model performance.

106. What is entropy in decision trees?

A measure of impurity in data.

107. What is information gain?

Reduction in entropy after a dataset split.

108. Name one algorithm to build decision trees.

ID3, C4.5, or CART.

109. What is pruning in decision trees?

Removing unnecessary branches to prevent overfitting.

110. What is Bayes theorem?

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

111. What is prior probability?

The probability of an event before evidence.

112. What is posterior probability?

Updated probability after considering evidence.

113. What is K in K-NN?

The number of nearest neighbors considered.

114. What distance measures are used in K-NN?

Euclidean, Manhattan, Minkowski.

115. What is bias-variance tradeoff?

Balance between underfitting and overfitting.

116. Define accuracy.

$$\frac{TP+TN}{TP+TN+FP+FN}$$

117. What is precision?

Proportion of predicted positives that are truly positive.

118. What is recall?

Proportion of actual positives correctly identified.